

# Large-Scale Machine Learning


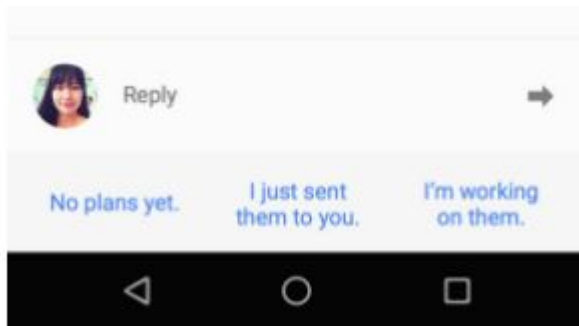
From DistBelief To TensorFlow

# Overview of this Presentation

- Review two large-scale machine learning systems from Google
  - DistBelief
  - TensorFlow



# Background: Massive Success of Deep Learning

- Various real-world applications
  - Image recognition
  - Speech recognition
  - Natural language processing
  - Game play (e.g., Go)
  - ...
- Beating existing machine-learning algorithms



# Challenges

- Require large amount of computational and storage resources
  - Increasing dataset size
  - Increasing number of neurons and connections
- Require rapid iteration on new learning algorithm development
  - Hot research topic that many people are actively working on
- Do not fit well with existing parallel/distributed programming models
  - E.g.) Mapreduce, Spark, graph execution engine, ...




# What Google Has Been Building

- DistBelief (2012)
  - Parallel and distributed execution of large-scale deep network
- TensorFlow (2015-)
  - More generalized data-flow execution engine
  - Still specialized for machine learning

**DistBelief**

# Before Diving Into...

- What is **deep learning**?
  - The modern reincarnation of Artificial Neural Networks from the 1980s and 90s
  - A collection of simple trainable mathematical units, which collaborates to compute a complicated function
  - Compatible with supervised, unsupervised, and reinforcement learning




# What is Neural Network?

- Approximate some function  $f^*$
- Typically represented by composing many different functions
  - $f(x) = f^{(3)}(f^{(2)}(f^{(1)}(x)))$


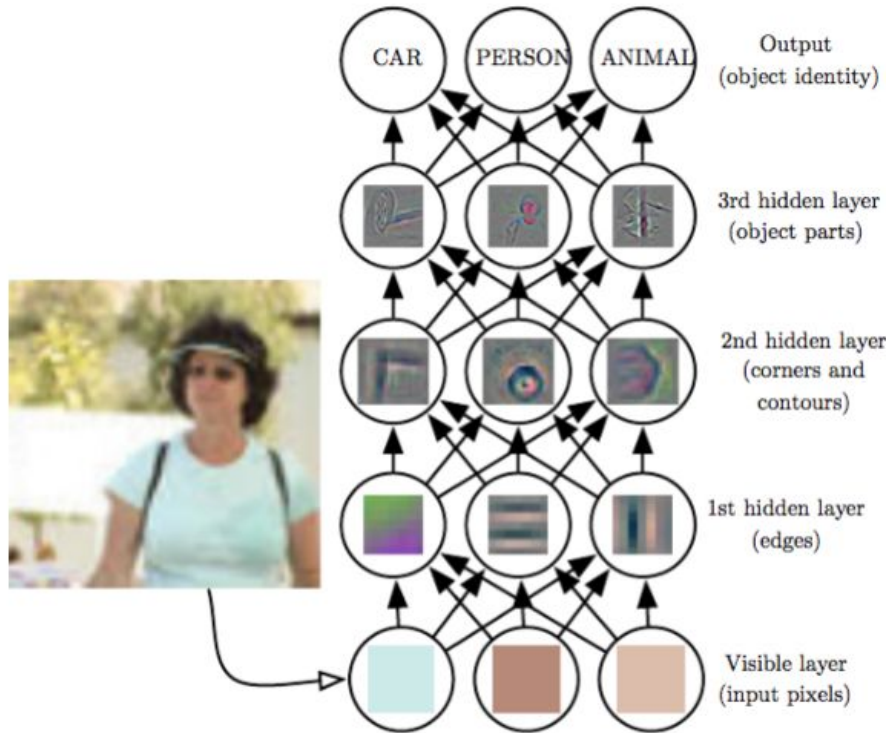
- Example: XOR

$$f(x) = [1 \quad -2] \max\{0, \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix} x + \begin{bmatrix} 0 \\ -1 \end{bmatrix}\}$$






# Example Deep Neural Networks



# Accuracy Improvement with Deeper Networks

Example: Transcribe multi-digit numbers from photographs of addresses



# How Can We Train a Neural Network?

**while** not done:


- pick a random training case ( $x$ ,  $y$ )

- run a neural network on input  $x$

- modify connection weights to make prediction closer to  $y$

# How to Modify Connections?

- Follow the gradient of the error w.r.t. the connection (e.g., weight parameters)



Gradient points in direction of improvement


# Compute Gradient with “Back-Propagation”

- Use the chain rule of calculus:

error rate  $\rightarrow$   $\frac{\partial z}{\partial x_i} = \sum_j \frac{\partial z}{\partial y_j} \frac{\partial y_j}{\partial x_i}$   $\leftarrow$  hidden node's parameter

learning parameter  $\rightarrow$

- Backtrack a network from output to input
- Memorize intermediate results to avoid recalculation




# Compute Gradient with “Back-Propagation”

- Use the chain rule of calculus:

error rate  $\longrightarrow$   $\frac{\partial z}{\partial x_i} = \sum_j \frac{\partial z}{\partial y_j} \frac{\partial y_j}{\partial x_i}$   $\longleftarrow$  hidden node's parameter

learning parameter  $\longrightarrow$

- Backtrack a network from output to input
- Memorize intermediate results to avoid recalculation



# Stochastic Gradient Descent (SGD) Learning

- Estimate the gradient with a **small set of samples** (minibatch)

**Require:** Learning rate  $\epsilon_k$ .

**Require:** Initial parameter  $\theta$

**while** stopping criterion not met **do**

Sample a minibatch of  $m$  examples from the training set  $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$  with corresponding targets  $\mathbf{y}^{(i)}$ .

Compute gradient estimate:  $\hat{\mathbf{g}} \leftarrow +\frac{1}{m} \nabla_{\theta} \sum_i L(f(\mathbf{x}^{(i)}; \theta), \mathbf{y}^{(i)})$

Apply update:  $\theta \leftarrow \theta - \epsilon \hat{\mathbf{g}}$

**end while**



Loss function


# How Can We Train Large Neural Nets Quickly?

- Exploit many kinds of parallelism
  - Model parallelism
  - Data parallelism
- Note: bad fit with MapReduce
  - Mutation to learning parameters
  - Non-deterministic result
  - “Weak” correctness guarantee
    - OK if we can train a model accurately
  - ...



# Model Parallelism

- Partition model across machines
  - The most densely connected areas are on the same partition
  - Up to 144 partitions with significant speedups




# Data Parallelism

- Distribute training ***across*** multiple model instances
- Propose two algorithms
  - Downpour SGD (= variant of asynchronous SGD)
  - Sandblaster L-GFGS


# Downpour SGD

- Divide the training data into a number of subsets
- Run a copy of the model on each of these subsets
- Communicate updates through a centralized parameter server



# Asynchronous Communication Between Servers

- Model replicas run independently of each other
  - Parameter server shards also run independently of one another
- 
- **Pros:** Can continue processing even when one machine is down
  - **Cons:** Additional stochasticity in the optimization procedure
    - E.g.) A model replica computes its gradients based on out-of-date parameters




# Optimizing Downpour SGD

- Fetch & pull parameters only every  $N$  steps
- “Warmstart” model training with only a single model replica before unleashing the other replicas
- Apply a non-fixed learning rate
  - “Adagrad” adaptive learning rate procedure

$$\eta_{i,K} = \gamma / \sqrt{\sum_{j=1}^K \Delta w_{i,j}^2}$$

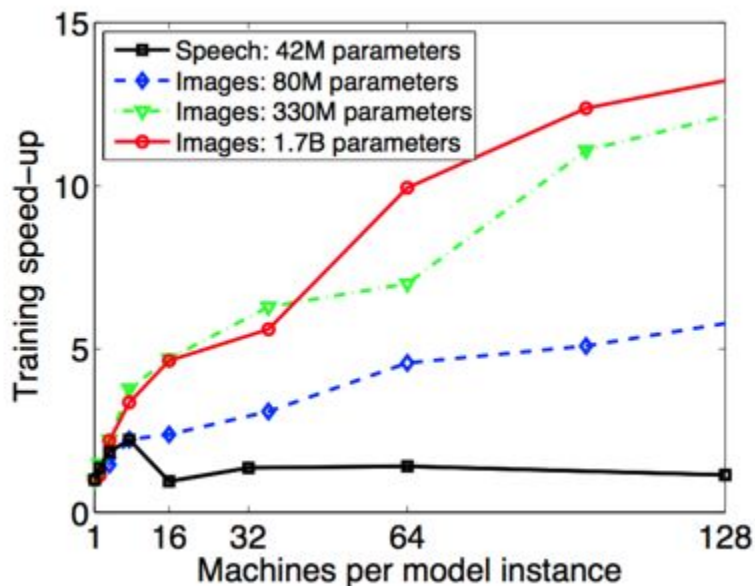
# Sandblaster L-BFGS

- Avoid high-frequent, high-bandwidth communication
  - Coordinator issues “commands” to Parameter Servers
    - Parameter Servers execute commands and store results
  - Coordinator dynamically assigns tasks to Model Replicas
    - Multiple copies of work can be scheduled to address a slow bottleneck machine





# Evaluation: Model Parallelism

- Mean time to process a min-batch for a simple SGD
- The largest model benefits the most
  - More than 12x speedup using 81 machines



# Evaluation: Data Parallelism

- Speech model in a variety of configurations





# Summary of DistBelief

- Parallel and distributed execution of deep learning with SGC
  - Partition model across machines
  - Distribute training across multiple model instances
- Individual techniques are not surprising, but achieved very good results

# TensorFlow

# Lesson Learned From DistBelief

- Need a better abstraction layer
  - E.g.) Allow users to add new primitives or without changing DistBelief core
- Need one system for both large-scale training and small-scale deployment
  - E.g.) Experiment a new algorithm on a single machine first and then use the same code for large-scale deployment
- Need to support heterogeneous hardware
  - GPU, custom ASIC (e.g., TensorFlow Processing Unit)

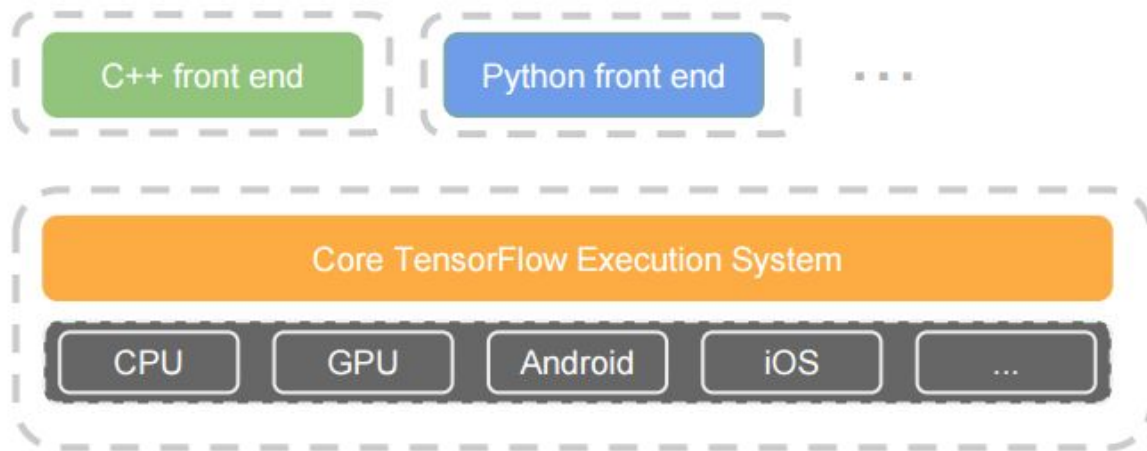
# TensorFlow™

- Second-generation system for the implementation and deployment of large-scale machine learning models
- Takes computations with a ***dataflow-like model*** and maps them onto a wide variety of different hardware platforms
  - Inference on mobile device platforms
  - Modest-sized training and inference on single machines with GPUs
  - Large-scale training on >100 specialized machines with >1000 GPUs



# Expressing High-Level ML Computations

- Core in C++
  - Very low overhead
- Different front ends for specifying/driving the computation
  - Python and C++ today, easy to add more



# Example TensorFlow Program (in Python)

```
import tensorflow as tf

b = tf.Variable(tf.zeros([100]))
W = tf.Variable(tf.random_uniform([784,100],-1,1))
x = tf.placeholder(name="x")
relu = tf.nn.relu(tf.matmul(W, x) + b)
C = [...]

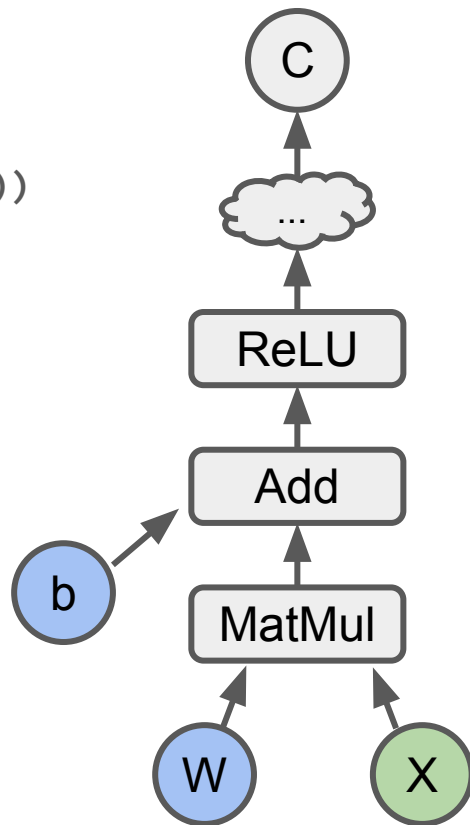
s = tf.Session()
for step in xrange(0, 10):
    input = ...construct 100-D input array ...
    result = s.run(C, feed_dict={x: input})
    print step, result
```

# Representation of Computation Graph

```
import tensorflow as tf

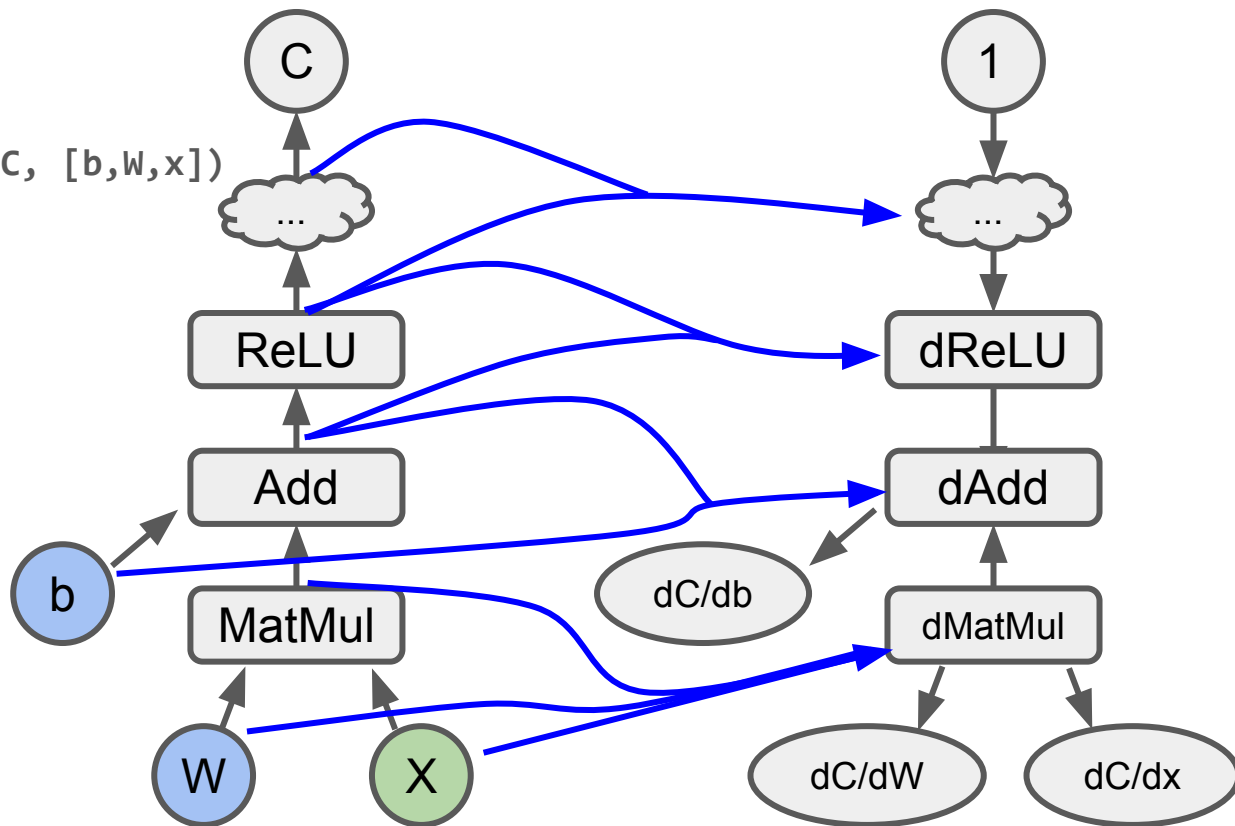
b = tf.Variable(tf.zeros([100]))
W = tf.Variable(tf.random_uniform([784,100],-1,1))
x = tf.placeholder(name="x")
relu = tf.nn.relu(tf.matmul(W, x) + b)
C = [...]

s = tf.Session()
for step in xrange(0, 10):
    input = ...construct 100-D input array ...
    result = s.run(C, feed_dict={x: input})
    print step, result
```



# Graph for Computing Gradients

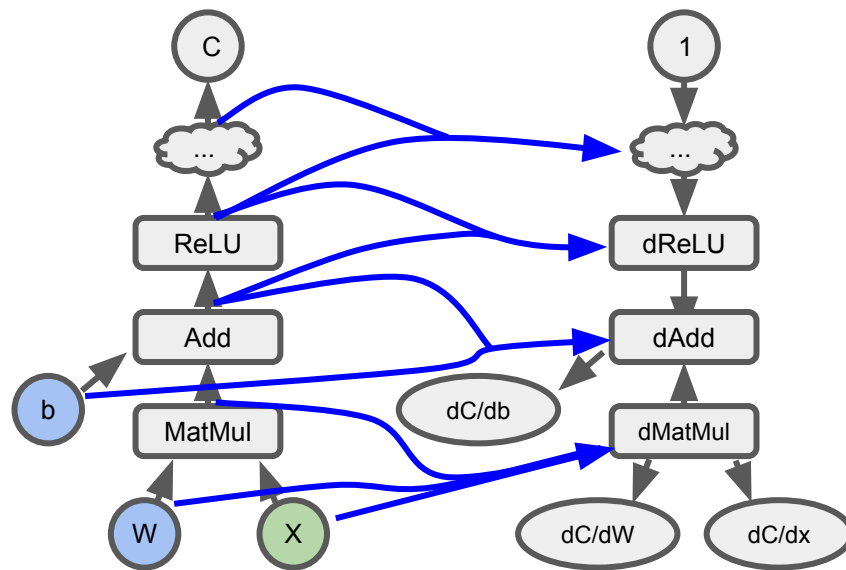
`[db,dW,dx] = tf.gradients(C, [b,W,x])`





# By Representing Gradient Computation as Graph..

- Allow Users to easily implement new algorithms for computing gradients
  - C.f.) Implementing the “momentum” algorithm in DistBelief required change to to the C++ parameter server and execution of arbitrary code in write operations




# TensorFlow Execution Model

- **Vertex (= operation)**

- An atomic unit of computation
- May have *mutable* state that can be shared between different executions of the graph
- E.g.) Add, MatMul, Variable, ReLU, Save, Merge, ...




- **Edge (= flow of tensors)**

- Input from or output to a vertex
- E.g.) Matrix multiplication takes two 2-D tensors and produces a 2-D tensor
- E.g.) Mini-batch 2-D convolution takes two 4-D tensors and produces another 4-D tensor




# Parallel and Concurrent Execution

- Client specifies a **subgraph** that should be executed
  - Zero or more edges to **feed** input tensors
  - One or more edges to **fetch** output tensors
- Multiple subgraphs are executed concurrently



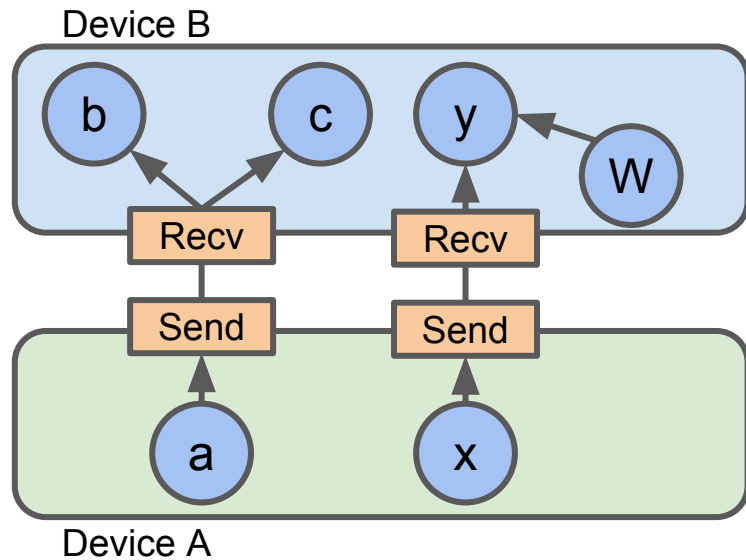


# Concurrent Execution in a Training Application

- **I/O subgraph:** Read records (e.g., images) from a distributed file system
- **Preprocessing subgraph:** Transform individual input records
  - E.g.) Decode images and apply random distortions
- **Training subgraph:** Update the model based on different input batches
  - Implement data-parallel training
  - Consist of many concurrent steps
- **Checkpointing subgraph:** Run periodically for fault tolerance



# Distributed Execution

1. Place operations on *feasible* devices (e.g., CPU, GPU)
  - E.g.) Only place this node on a device of type GPU
2. Partition operations into per-device subgraphs
  - Per-device subgraph contains all of the operations assigned to the device, with Send and Recv

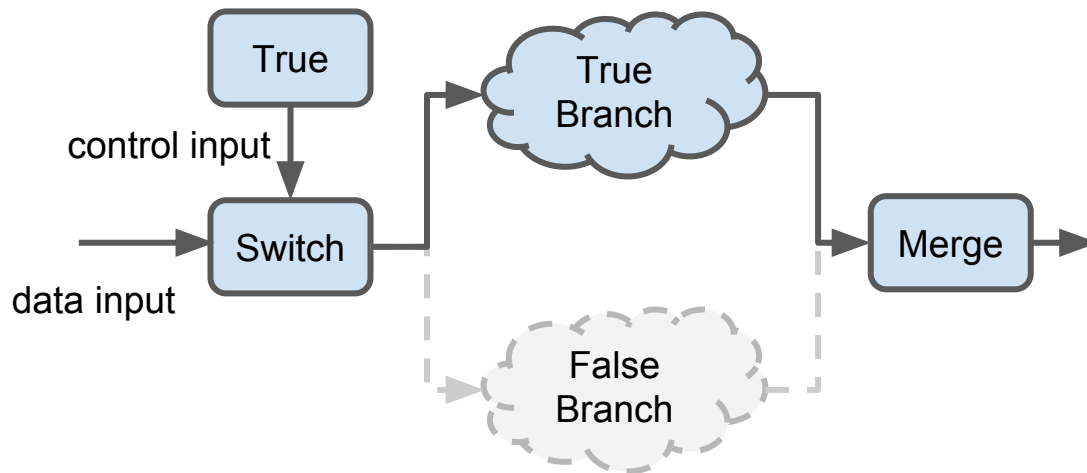


# Dynamic Control Flow

- Most evaluation in TensorFlow is ***strict***
  - All inputs to an operation must be computed before the operation executes
- However, some advanced algorithms requires ***non-strict*** evaluation
  - E.g.) efficient training of a recurrent neural network

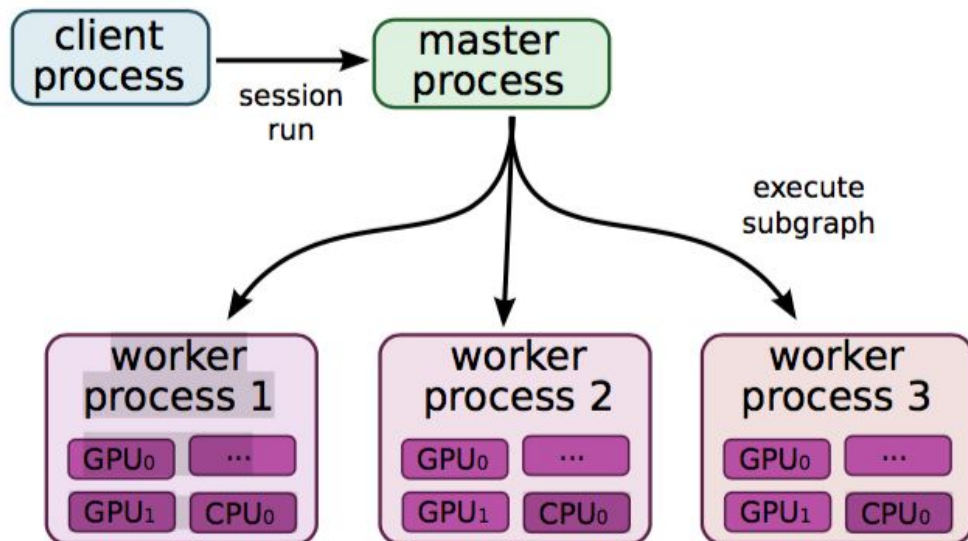
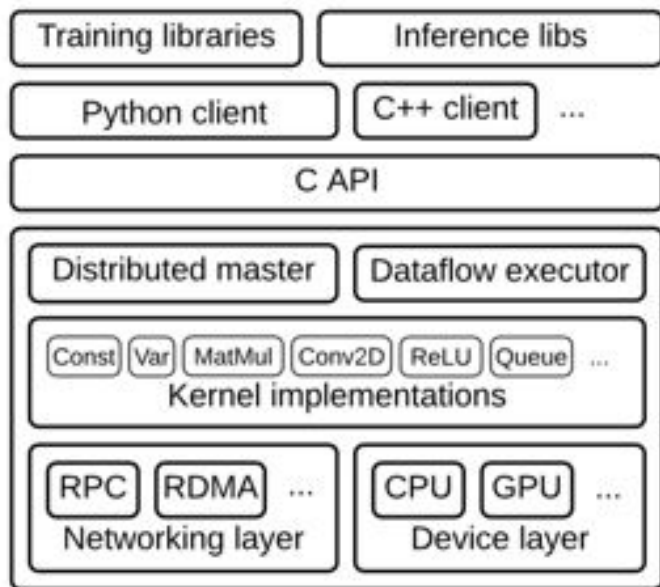
# Dynamic Control Flow

- Introduce Switch and Merge
- Can also express iteration



# Implementation Sketch

- Master-worker architecture





# Optimizations & Challenges

- Typical compiler optimizations
  - Common subexpression elimination
  - Fusion (e.g., replace multiple loops with a single one)
- Lossy compression
  - Convert a 32-bit float to a 16-bit float when sending data between devices
  - Not IEEE 16-bit floating point standard, but with 16 bits less precision in the significand
- Synchronous replica coordination
- Node placement and scheduling
- ....


$$1.2345 = \underbrace{12345}_{\text{significand}} \times \underbrace{10^{-4}}_{\text{base}}^{\text{exponent}}$$

# Synchronous Replica Coordinations

- Is **synchronous** training really a bad idea?
  - GPUs enable training with hundreds (not thousands) of machines




- Synchronous replication with **backup workers**
  - Improve throughput by up to 15% by addressing stragglers



# Basic Node Placement Algorithm

1. Compute the feasible set of devices for each node
2. Compute the graph components that must be placed together
  - Use union-find on the graph of colocation constraints
3. Compute the intersection of the feasible device sets for each component
4. Run the placement simulator



# Greedy Heuristic Placement Simulator

- Examines the completion time of a node on each possible device
  - Estimated (or measured) execution time of the operation
  - Communication cost for transmitting inputs to the node
- Selects the device where the node's operation would finish soonest

# Optimizations of Node Placement and Scheduling

- Scheduling of RECV nodes for reading remote values
  - Estimate when to start the RECV nodes by analyzing the critical paths
  - Perform as-soon-as-possible/as-late-as-possible (ASAP/ALAP) calculation
- Memory management for “back-propagation” gradient calculation
  - Manage GPU memory when iterating over long sequences in the input data
  - E.g.) keeping intermediate data v.s. recomputing
- ...
- ...

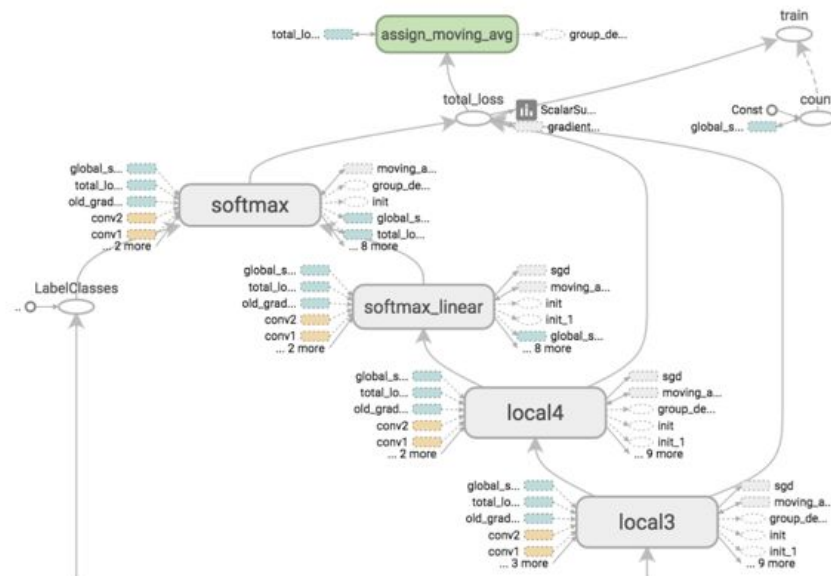
# Note on Engineering Efforts

- Great engineering accomplishment!

- Complex mathematical operations
- Heterogeneous environments
- Stochastic behavior
- Parallel and distributed execution
- ...

- Various supporting tools developed

- Graph visualization
- Performance visualization




# Evaluation: Single Machine Benchmarks

- Comparable performance
  - TensorFlow and Torch use the same version of the cuDNN library
  - The Neon library uses hand-optimized convolutional kernels

Library	Training step time (ms)			
	AlexNet	Overfeat	OxfordNet	GoogleNet
Caffe [36]	324	823	1068	1935
Neon [56]	87	<b>211</b>	<b>320</b>	<b>270</b>
Torch [17]	<b>81</b>	268	529	470
TensorFlow	<b>81</b>	279	540	445

# Evaluation: Image Classification

- Investigate the scalability of training Google's Inception v-3 model
  - 17 Parameter Server tasks and a varying number of worker tasks
  - Each Parameter Server task has 8 IvyBridge cores
  - Each worker task has 5 IvyBridge and one NVIDIA K40 GPU
- Throughput improves to 2,300 images per second (with 200 workers)





# Summary

- Two large-scale machine learning systems
  - DistBelief
  - TensorFlow
- TensorFlow is available at <https://www.tensorflow.org/>
  - Used by various researchers and companies (e.g., UBER, Snapchat, ARM, Airbus)

# References

- <http://research.google.com/pubs/jeff.html>
  - TensorFlow: A system for large-scale machine learning (2016)
  - TensorFlow: Large-Scale Machine Learning on Heterogeneous Distributed Systems (2015)
  - Large Scale Distributed Deep Networks (2012)
  - ...
- <http://www.deeplearningbook.org/>